

Parallel Computation for the All-Pairs Suffix-Prefix Problem

Felipe A. Louza¹ Simon Gog² Leandro Zanolto¹
Guido Araujo¹ Guilherme P. Telles¹

¹Institute of Computing, UNICAMP, Brazil

²Institute of Theoretical Informatics, KIT, Germany

SPIRE'16
Beppu, Japan

Outline

1. Introduction

2. Preliminaries

3. Related work

4. Parallel Algorithm

5. Experiments

6. Conclusion

Introduction

All-pairs suffix-prefix matching problem (APSP):

- ▶ Given a collection of strings $\mathcal{S} = S^1, S^2, \dots, S^m$, and a threshold τ .
- ▶ APSP is to find, for all pairs S^i and S^j , the longest suffix of S^i that overlaps S^j that is larger than τ .
- ▶ Suffix-prefix match (overlap):




Figure: S^i overlaps S^j , the suffix $S^i[a, n_i - 1]$ is equal to the prefix $S^j[1, b]$

Motivation:

- ▶ DNA assembly (bottleneck stage).
- ▶ EST clustering.
- ▶ Approximating the shortest common superstring.

Introduction

Practical Algorithms:

- ▶ DNA assemblers provide fast solutions with non-optimal algorithms.
- ▶ SGA [Simpson and Durbin, 2010] and Readjoiner [Gonnella and Kurtz, 2012] have isolated overlap detection stage.
- ▶ [Rachid and Malluhi, 2015] presented SOF to solve the APSP for DNA sequences. SOF has a good performance with multiple threads¹.

Optimal Algorithms:

- ▶ [Gusfield et al., 1992] solved the APSP in optimal time → suffix trees and stacks.
- ▶ [Ohlebusch and Gog, 2010] improved memory usage and running time → enhanced suffix arrays and stacks.
- ▶ [Tustumi et al., 2016] proposed a different traversal of the enhanced suffix array and replaced stacks by linked lists to achieve a better practical running time.

¹SGA, Readjoiner and SOF may be executed in multithreading environments.

Introduction

Our proposal:

- ▶ We showed how to **parallelize** the optimal algorithm by [Tustumi et al., 2016] to solve the APSP.
- ▶ Our parallel algorithm achieves a **consistent speedup** with a **small memory footprint** when compared with [Tustumi et al., 2016].
- ▶ Also, it is competitive with **SOF** when the threshold τ is small.

Workflow:

- ▶ Our algorithm is composed by three phases.
- ▶ We **separated the computation** of the **local solution** followed by the **global solution** and by the **identical suffixes** searching into independent phases.

Outline

1. Introduction

2. Preliminaries

3. Related work

4. Parallel Algorithm

5. Experiments

6. Conclusion

Preliminaries:

Let $S[1, n]$ be a string of length $|S| = n$ over an ordered alphabet Σ .

- ▶ A prefix is a substring of the form $S[1, i]$.
- ▶ A suffix is a substring $S[i, n]$ that will be denoted by S_i .

The **suffix array** of $S[1, n]$, SA, is an array of integers in $[1, n]$ that gives the lexicographic order of all suffixes:

- ▶ $S_{SA[1]} < S_{SA[2]} < \dots < S_{SA[n]}$;
- ▶ We denote the **position of S_i** in SA as $\text{pos}(S_i)$.

The **LCP-array** stores the length of the longest common prefix (lcp) of two consecutive suffixes in SA:

- ▶ $\text{LCP}[i] = lcp(S_{SA[i]}, S_{SA[i-1]})$ for $1 < i \leq n$, and
- ▶ $\text{LCP}[1] = 0$.

The **range minimum query** (rmq) on LCP gives the smallest lcp value in an interval of SA:

- ▶ $rmq(i, j) = \min_{i < k \leq j} \{\text{LCP}[k]\}$.
- ▶ It is well known that $lcp(S_{SA[i]}, S_{SA[j]}) = rmq(i, j)$, with $1 \leq i < j \leq n$.

Preliminaries:

Let $\mathcal{S} = S^1, S^2, \dots, S^m$ be a **collection of strings** of lengths $n_i = |S^i|$.

- ▶ $S^{cat} = S^1 \cdot \$_1 \cdot S^2 \cdot \$_2 \cdots S^m \cdot \$_m$ is the concatenated string of length $N = m + \sum_{i=1}^m n_i$.
- ▶ Each $\$_i$ is a **distinct separator** not in Σ that precedes $\forall \alpha \in \Sigma$, and $\$_i < \$_j$ if $i < j$.
- ▶ $S_k^\$$ denotes the prefix of S_k^{cat} that ends at the first separator $\$_j$.

The **generalized suffix array** of \mathcal{S} , GSA, is the SA of the concatenated string S^{cat} .

For a clearer notation, we introduce:

- ▶ STR indicates which string in \mathcal{S} a suffix came from, $STR[i] = j$ if $S_{SA[i]}^\$$ ends with $\$_j$.
- ▶ SA' holds the position of a suffix with respect to the string it came from (up to the separator), defined as $SA'[i] = k$ if $S_{SA[i]}^\$ = S_k^j \$_j$.

GESA denotes the GSA enhanced with the arrays STR, SA' and LCP.

Preliminaries:

Let S^k be the j -th (lexicographically) **smallest string** in \mathcal{S} .

- ▶ P is an array of $m + 1$ **integers** that stores in $P[j] = \text{pos}(S^k[1, n_k])$.
- ▶ We define $P[0] = m + 1$.

Let $B^j = (P[j - 1], P[j])^2$ be a **block** of GESA corresponding to S^k .

- ▶ GESA can be partitioned into m blocks B^1, B^2, \dots, B^m , one for each string S^k in \mathcal{S} .

	i	SA	LCP	STR	SA'	$S_{SA[i]}^s$
P[0] →	1	4	0	1	4	$\$1$
	2	8	0	2	4	$\$2$
	3	11	0	3	3	$\$3$
	4	15	0	4	4	$\$4$
	5	7	0	2	3	a $\$2$
P[1] →	6	10	1	3	2	a $\$3$
	7	14	1	4	3	a $\$4$
	8	9	1	3	1	aa $\$3$
P[2] →	9	13	2	4	2	aa $\$4$
	10	1	2	1	1	aac $\$1$
P[3] →	11	2	1	1	2	ac $\$1$
	12	5	2	2	1	aca $\$2$
	13	3	0	1	3	c $\$1$
P[4] →	14	6	1	2	2	ca $\$2$
	15	12	2	4	1	caa $\$4$

Figure: GESA of $\mathcal{S} = \{\text{aac}, \text{aca}, \text{aa}, \text{caa}\}$. Suffixes in block 1 are highlighted.

²open interval: $P[j - 1] + 1$ stores the $\text{lcp}(S_{SA[P[j-1]+1]}, S_{SA[P[j]-1]})$

Outline

1. Introduction

2. Preliminaries

3. Related work

4. Parallel Algorithm

5. Experiments

6. Conclusion

Related work:

The algorithm by [Tustumi et al., 2016] solves the APSP in optimal $O(N + m^2)$ time, based on the following remarks:

- ▶ All suffixes that overlap S^k are in positions prior to $\text{pos}(S^k)$ or are identical to S^k and directly succeed $\text{pos}(S^k)$.
- ▶ If two different suffixes of S^r are a prefix of S^k , the longest is closer to $\text{pos}(S^k)$.
- ▶ Given two prefixes $S^t < S^k$, if a suffix of S^r , of length ℓ , overlap S^t and $\ell > \text{lcp}(S^t, S^k)$, then such suffix does not overlap S^k .

	i	SA	LCP	STR	SA'	$S_{SA[i]}^s$
	1	4	0	1	4	$\$1$
	2	8	0	2	4	$\$2$
	3	11	0	3	3	$\$3$
	4	15	0	4	4	$\$4$
P[0] →	5	7	0	2	3	a $\$2$
	6	10	1	3	2	a $\$3$
	7	14	1	4	3	a $\$4$
P[1] →	8	9	1	3	1	aa $\$3$
	9	13	2	4	2	aa $\$4$
P[2] →	10	1	2	1	1	aac $\$1$
	11	2	1	1	2	ac $\$1$
P[3] →	12	5	2	2	1	aca $\$2$
	13	3	0	1	3	c $\$1$
	14	6	1	2	2	ca $\$2$
P[4] →	15	12	2	4	1	caa $\$4$

Figure: GESA of $S = \{aac, aca, aa, caa\}$.

Related work:

The algorithm by [Tustumi et al., 2016] solves the APSP in optimal $O(N + m^2)$ time, based on the following remarks:

- ▶ All suffixes that overlap S^k are in positions prior to $\text{pos}(S^k)$ or are identical to S^k and directly succeed $\text{pos}(S^k)$.
- ▶ If two different suffixes of S^r are a prefix of S^k , the longest is closer to $\text{pos}(S^k)$.
- ▶ Given two prefixes $S^t < S^k$, if a suffix of S^r , of length ℓ , overlap S^t and $\ell > \text{lcp}(S^t, S^k)$, then such suffix does not overlap S^k .

	i	SA	LCP	STR	SA'	$S_{SA[i]}^s$
P[0] →	1	4	0	1	4	$\$1$
	2	8	0	2	4	$\$2$
	3	11	0	3	3	$\$3$
	4	15	0	4	4	$\$4$
	5	7	0	2	3	a $\$2$
	6	10	1	3	2	a $\$3$
P[1] →	7	14	1	4	3	a $\$4$
	8	9	1	3	1	aa $\$3$
P[2] →	9	13	2	4	2	aa $\$4$
	10	1	2	1	1	aac $\$1$
P[3] →	11	2	1	1	2	ac $\$1$
	12	5	2	2	1	aca $\$2$
	13	3	0	1	3	c $\$1$
	14	6	1	2	2	ca $\$2$
P[4] →	15	12	2	4	1	caa $\$4$

Figure: GESA of $\mathcal{S} = \{\text{aac}, \text{aca}, \text{aa}, \text{caa}\}$.

Related work:

The algorithm by [Tustumi et al., 2016] solves the APSP in optimal $O(N + m^2)$ time, based on the following remarks:

- ▶ All suffixes that overlap S^k are in positions prior to $\text{pos}(S^k)$ or are identical to S^k and directly succeed $\text{pos}(S^k)$.
- ▶ If two different suffixes of S^r are a prefix of S^k , the longest is closer to $\text{pos}(S^k)$.
- ▶ Given two prefixes $S^t < S^k$, if a suffix of S^r , of length ℓ , overlap S^t and $\ell > \text{lcp}(S^t, S^k)$, then such suffix does not overlap S^k .

	i	SA	LCP	STR	SA'	$S_{SA[i]}^s$
P[0] →	1	4	0	1	4	$\$1$
	2	8	0	2	4	$\$2$
	3	11	0	3	3	$\$3$
	4	15	0	4	4	$\$4$
	5	7	0	2	3	$a\$2$
	6	10	1	3	2	$a\$3$
	7	14	1	4	3	$a\$4$
P[1] →	8	9	1	3	1	$aa\$3$
	9	13	2	4	2	$aa\$4$
P[2] →	10	1	2	1	1	$aac\$1$
	11	2	1	1	2	$ac\$1$
P[3] →	12	5	2	2	1	$aca\$2$
	13	3	0	1	3	$c\$1$
P[4] →	14	6	1	2	2	$ca\$2$
	15	12	2	4	1	$caa\$4$

Figure: GESA of $S = \{aac, aca, aa, caa\}$.

Related work:

Algorithm [Tustumi et al., 2016]:

- ▶ The blocks are processed in order, B^1, B^2, \dots, B^m .
- ▶ For each B^j : suppose that $\text{pos}(S^k) = P[j]$ and $\text{pos}(S^t) = P[j - 1]$.
 1. **Local solution** is found scanning B^j backwards;
 2. **Global solution** is obtained reusing the local solutions of the previous blocks.
 3. **Identical suffixes** are found scanning GESA forward from $P[j] + 1$.
- ▶ The algorithm uses m local lists and m global lists to track all overlaps seen so far.

	i	SA	LCP	STR	SA'	$S_{SA[i]}^i$
	1	4	0	1	4	$\$1$
	2	8	0	2	4	$\$2$
	3	11	0	3	3	$\$3$
	4	15	0	4	4	$\$4$
$P[0] \rightarrow$	5	7	0	2	3	$a\$2$
	6	10	1	3	2	$a\$3$
	7	14	1	4	3	$a\$4$
$P[1] \rightarrow$	8	9	1	3	1	$aa\$3$
	9	13	2	4	2	$aa\$4$
$P[j-1] \rightarrow$	10	1	2	1	1	$aac\$1 \leftarrow S^t$
	11	2	1	1	2	$ac\$1$
$P[j] \rightarrow$	12	5	2	2	1	$aca\$2 \leftarrow S^k$
	13	3	0	1	3	$c\$1$
	14	6	1	2	2	$ca\$2$
$P[4] \rightarrow$	15	12	2	4	1	$caa\$4$

Figure: GESA of $\mathcal{S} = \{aac, aca, aa, caa\}$.

Related work:

Algorithm [Tustumi et al., 2016]:

► Local solution:

- For block B^j :
- GESA is scanned backwards, from $P[j]$ to $P[j - 1] + 1$.
- $\ell = rmq(i, P[j])$ is computed in $O(1)$ time during the scanning of B^j .
 - if $|S_{SA'[i]}^{STR[i]}| = \ell$ then $S^{STR[i]}$ overlaps S^k in ℓ symbols, and $insert_at_end(L_{local}[STR[i]], \ell)$.
 - At the end, the longest overlaps in B^j are at the front of each local list.

	i	SA	LCP	STR	SA'	$S_{SA[i]}^S$
	1	4	0	1	4	S_1
	2	8	0	2	4	S_2
	3	11	0	3	3	S_3
	4	15	0	4	4	S_4
$P[0] \rightarrow$	5	7	0	2	3	aaS_2
	6	10	1	3	2	aS_3
$P[1] \rightarrow$	7	14	1	4	3	aaS_4
	8	9	1	3	1	aaS_3
	9	13	2	4	2	aaS_4
$P[2] \rightarrow$	10	1	2	1	1	$aacS_1$
	11	2	1	1	2	acS_1
$P[3] \rightarrow$	12	5	2	2	1	$acaS_2$
	13	3	0	1	3	cS_1
	14	6	1	2	2	caS_2
$P[4] \rightarrow$	15	12	2	4	1	$caaS_4$

i	$L_{local}[1]$	$L_{local}[2]$	$L_{local}[3]$	$L_{local}[4]$
8	[]	[]	[]	[]
7	[]	[]	[]	[1]
6	[]	[]	[1]	[1]
5	[]	[1]	[1]	[1]

Figure: GESA of $S = \{aac, aca, aa, caa\}$.

Related work:

Algorithm [Tustumi et al., 2016]:

► Global solution:

- The longest suffix of S^r that overlaps S^k may be positioned in a previous block.
- The global lists store these overlaps.
- Each $L_{global}[r]$ is updated as each block is processed.
 - First, we remove suffixes larger than $lcp(S^t, S^k)$ from local lists.
 - $L_{local}[r]$ is prepended in the global list $L_{global}[r]$.
 - The first element of each $L_{global}[r]$ is the longest overlap of S^r to S^k , that is $Ov[k, r]$.

	i	SA	LCP	STR	SA'	$S_{SA[i]}^s$
	1	4	0	1	4	$\$1$
	2	8	0	2	4	$\$2$
	3	11	0	3	3	$\$3$
	4	15	0	4	4	$\$4$
$P[0] \rightarrow$	5	7	0	2	3	$a\$2$
	6	10	1	3	2	$a\$3$
	7	14	1	4	3	$a\$4$
$P[1] \rightarrow$	8	9	1	3	1	$aa\$3$
	9	13	2	4	2	$aa\$4$
$P[2] \rightarrow$	10	1	2	1	1	$aac\$1$
	11	2	1	1	2	$aca\$1$
$P[3] \rightarrow$	12	5	2	2	1	$aca\$2$
	13	3	0	1	3	$c\$1$
	14	6	1	2	2	$ca\$2$
$P[4] \rightarrow$	15	12	2	4	1	$caa\$4$

i	$L_{local}[1]$	$L_{local}[2]$	$L_{local}[3]$	$L_{local}[4]$
12	[]	[]	[]	[]
11	[2]	[]	[]	[]
10	[2]	[]	[]	[]

$P[j-1]$	$L_{global}[1]$	$L_{global}[2]$	$L_{global}[3]$	$L_{global}[4]$
5	[]	[1]	[1]	[1]
8	[]	[1]	[2,1]	[2,1]
10	[2]	[1]	[1]	[1]

Figure: GESA of $S = \{aac, aca, aa, caa\}$.

Related work:

Algorithm [Tustumi et al., 2016]:

► Identical suffixes:

- We scan GESA forward from $i = P[j] + 1$ to q , while $LCP[q] = n_k$.
- The length of these suffixes are inserted in Ov , possibly overwriting the results in $L_{global}[r]$, which is correct as such overlaps are larger.

	i	SA	LCP	STR	SA'	$S_{SA[i]}^s$
	1	4	0	1	4	$\$1$
	2	8	0	2	4	$\$2$
	3	11	0	3	3	$\$3$
	4	15	0	4	4	$\$4$
P[0] →	5	7	0	2	3	$aa\$2$
	6	10	1	3	2	$aa\$3$
	7	14	1	4	3	$aa\$4$
P[1] →	8	9	1	3	1	$aa\$3$
	9	13	2	4	2	$aa\$4$
P[2] →	10	1	2	1	1	$aac\$1$
	11	2	1	1	2	$ac\$1$
P[3] →	12	5	2	2	1	$aca\$2$
	13	3	0	1	3	$c\$1$
	14	6	1	2	2	$ca\$2$
P[4] →	15	12	2	4	1	$caa\$4$

$$Ov[m, m]^3 =$$

	1	2	3	4
1		1	2	2
2	2		1	1
3		1		2
4	1	2		

Figure: GESA of $S = \{aac, aca, aa, caa\}$.

³ $Ov[k, r]$ the length of the longest suffix of S^r that overlaps S^k .

Outline

1. Introduction
2. Preliminaries
3. Related work
4. Parallel Algorithm
5. Experiments
6. Conclusion

Parallel Algorithm:

We **split the computation** of all overlaps to solve the APSP in parallel:

At a glance, our algorithm is composed by:

1. Compute all **local solutions** scanning the blocks B^j concurrently.
2. Compute all **global solutions** accessing the local lists in parallel.
3. Compute all **identical suffixes** of all strings scanning the GESA in parallel.

Algorithm: Suppose that for each block B^j , $\text{pos}(S^k) = P[j]$ and $\text{pos}(S^t) = P[j - 1]$.

	i	SA	LCP	STR	SA'	$S_{SA[i]}^i$
P[0] →	1	4	0	1	4	S_1^1
	2	8	0	2	4	S_2^2
	3	11	0	3	3	S_3^3
	4	15	0	4	4	S_4^4
	5	7	0	2	3	aS_2^2
P[1] →	6	10	1	3	2	aS_3^3
	7	14	1	4	3	aS_4^4
	8	9	1	3	1	aaS_3^3
P[j-1] →	9	13	2	4	2	aaS_4^4
	10	1	2	1	1	$aacS_1^1 \leftarrow S^t$
P[j] →	11	2	1	1	2	acS_1^1
	12	5	2	2	1	$acaS_2^2 \leftarrow S^k$
P[4] →	13	3	0	1	3	cS_1^1
	14	6	1	2	2	caS_2^2
	15	12	2	4	1	$caaS_4^4$


Figure: GESA of $\mathcal{S} = \{aac, aca, aa, caa\}$.

Parallel Algorithm:

1. Local solutions:

- ▶ Scan all the m blocks **in parallel**.
 - ▶ Each block B^j is scanned backwards.
 - ▶ Whenever $|S_{SA'[i]}^{r=STR[j]}| = rmq(P[j], i)^4$ then a suffix of S^r overlaps S^k .
- ▶ We store the overlaps at a squared matrix of local lists:
 - ▶ $L_{local}[r][j]$ stores the length of a suffix of S^r that overlaps S^k .
 - ▶ The overlaps are ordered decreasingly by their lengths due to the backward scan.
- ▶ We compute $Min[j] = rmq(P[j - 1], P[j]) = lcp(S^k, S^t)$
- ▶ At the end, all local solutions have been computed and are stored into the local lists.

	i	SA	LCP	STR	SA'	$S_{SA[i]}^s$
	1	4	0	1	4	$\$1$
	2	8	0	2	4	$\$2$
	3	11	0	3	3	$\$3$
	4	15	0	4	4	$\$4$
P[0] →	5	7	0	2	3	$a\$2$
	6	10	1	3	2	$a\$3$
	7	14	1	4	3	$a\$4$
P[1] →	8	9	1	3	1	$aa\$3$
	9	13	2	4	2	$aa\$4$
P[j-1] →	10	1	2	1	1	$aac\$1 \leftarrow S^t$
	11	2	1	1	2	$ac\$1$
P[j] →	12	5	2	2	1	$aca\$2 \leftarrow S^k$
	13	3	0	1	3	$c\$1$
	14	6	1	2	2	$ca\$2$
P[4] →	15	12	2	4	1	$caa\$4$




⁴The $rmqs$ are solved as LCP is scanned, ℓ starts with ∞ and $\ell = \min(\ell, LCP[i + 1])$;

Parallel Algorithm:

2. Global solutions:

- ▶ We process the m lines of matrix L_{local} **in parallel**.
 - ▶ Each local list $L_{local}[r][j]$ corresponds to the suffixes of S^r that overlap S^k in B^j .
- ▶ For each line r (in parallel):
 - ▶ We process lists $L_{local}[r][j]$ in order $j = 1, 2, \dots, m$.
 - ▶ We use a **global list** L_{global} that is updated according to $\text{Min}[j] = \text{lcp}(S^t, S^k)$.
 - ▶ L_{global} is initially empty and for each $j = 1, 2, \dots, m$, all suffixes larger than $\text{Min}[j]$ are removed, these suffixes are **no longer overlaps** for the next blocks.
 - ▶ $L_{local}[r][j]$ is prepended to L_{global}
 - ▶ $\text{Ov}[r, k] \leftarrow \text{first}(L_{global})^5$.

	i	SA	LCP	STR	SA'	$S_{SA[i]}^s$
	1	4	0	1	4	$\$1$
	2	8	0	2	4	$\$2$
	3	11	0	3	3	$\$3$
	4	15	0	4	4	$\$4$
P[0] →	5	7	0	2	3	$a\$2$
	6	10	1	3	2	$a\$3$
	7	14	1	4	3	$a\$4$
P[1] →	8	9	1	3	1	$aa\$3$
	9	13	2	4	2	$aa\$4$
P[j-1] →	10	1	2	1	1	$aac\$1 \leftarrow S^t$
	11	2	1	1	2	$ac\$1$
P[j] →	12	5	2	2	1	$aca\$2 \leftarrow S^k$
	13	3	0	1	3	$c\$1$
	14	6	1	2	2	$ca\$2$
P[4] →	15	12	2	4	1	$caa\$4$



Local lists, lines are being accessed in parallel

⁵The longest suffix of S^r that overlaps S^k will be the first element of L_{global} .

Parallel Algorithm:

3. Identical suffixes:

- ▶ Scan all **the end of blocks B^j** concurrently.
- ▶ For each block B^j and its corresponding string S^k :
 - ▶ All suffixes of S^r identical to S^k appear directly after $\text{pos}(S^k)$.
 - ▶ We scan from $P[j] + 1$ while the next suffixes have the same length of S^k .
 - ▶ $\text{Ov}[r, k] \leftarrow \text{LCP}[i]^6$.

	i	SA	LCP	STR	SA'	$S_{SA[i]}^s$
P[0] →	1	4	0	1	4	$\$1$
	2	8	0	2	4	$\$2$
	3	11	0	3	3	$\$3$
	4	15	0	4	4	$\$4$
	5	7	0	2	3	$a\$2$
P[1] →	6	10	1	3	2	$a\$3$
	7	14	1	4	3	$a\$4$
	8	9	1	3	1	$aa\$3$
P[2] →	9	13	2	4	2	$aa\$4$
P[2] →	10	1	2	1	1	$aac\$1$
P[3] →	11	2	1	1	2	$ac\$1$
	12	5	2	2	1	$aca\$2$
	13	3	0	1	3	$c\$1$
P[4] →	14	6	1	2	2	$ca\$2$
	15	12	2	4	1	$caa\$4$

⁶Different from [Tustumi et al., 2016] $\text{Ov}[r, k]$ the length of the longest suffix of S^r that overlaps S^k .

Parallel Algorithm:

Theoretical costs:

► Time:

- Our algorithm runs in $O(N + m^2/t)$ time, where t is the number of threads.
 - Local solution: $O(\max_{1 \leq j \leq m} |B^j|)$ to scan all blocks.
 - Global solution: $O(m^2/t)$ to access all local lists.
 - Identical suffixes: $O(N)$, since it reads at most N elements of GESA.
- The worst case happens only when the string lengths are very unbalanced.
- In realistic cases ($m \gg t$ and all strings with about the same size) the parallel time is close to $N/t + m^2/t$.

► Space:

- Our algorithm uses $O(N + m^2)$ of space.
 - which is equal to the sequential algorithm by [Tustumi et al., 2016].
- The space is given by the $O(N)$ space of the GESA, and by the $O(m^2 + N)$ space of the matrix of local lists, where each list stores at most N overlaps.

Outline

1. Introduction
2. Preliminaries
3. Related work
4. Parallel Algorithm
- 5. Experiments**
6. Conclusion

Experiments:

Implementation:

- ▶ C++ using OpenMP.
- ▶ sdfs-lite v.2 to construct the GESA.
- ▶ Source code: <https://github.com/felipelouza/p-apsp>.

Comparison:

- ▶ **p-apsp**: our parallel algorithm.
- ▶ **apsp**: sequential optimal algorithm by [Tustumi et al., 2016].
- ▶ **SOF**: practical (non-optimal) solution by [Rachid and Malluhi, 2015].

Number of threads:

- ▶ $t = \{1, 2, 4, 8, 16, 32\}$
- ▶ set by `omp_set_num_threads()` for **p-apsp** and **SOF**

Configuration:

- ▶ 64 bits Debian GNU/Linux 8 (kernel 3.16.0-4)
- ▶ Intel Xeon Processor E5-2630 v3 20M Cache 2.40-GH
- ▶ **384 GB** of RAM.

Experiments:

Dataset:

- ▶ We used $m = 300.000$ strings from [EST database of *C. elegans*⁷](#).

Minimum overlap length (τ):

- ▶ We limited the number of overlaps by $\tau = \{5, 10, 15, 20\}$

Table: Number of overlaps found with 100.000, 200.000 and 300.000 ESTs varying τ .

m/τ	5	10	15	20
100.000	18,853,491	206,154	88,725	82,427
200.000	71,451,170	2,675,759	2,139,431	2,077,125
300.000	162,135,112	7,044,274	5,800,397	5,617,779

Observation:

- ▶ The number of overlaps decreases as τ increases.
- ▶ For 300.000, when $\tau = 5$ the number of overlaps is **23 times** larger than when $\tau = 10$.
- ▶ Such variation impacts the performance of all algorithms.

⁷<http://www.uni-ulm.de/in/theo/research/seqana.html>

Experiments:

Running Time⁸:

- ▶ SOF was the **fastest** in all experiments.
- ▶ **p-apsp** has shown a good performance for small threshold τ .
- ▶ There is an overhead with $\tau = 5$ when comparing **p-apsp₁** to **apsp**.




Figure: Running time for varying values of τ (`omp_get_wtime()`).

⁸not accounting for the time to build the GESA for **apsp** and **p-apsp**, and compact prefix tree for **SOF**.

Experiments:

Speedup:

- ▶ We evaluated the speedup of **p-apsp** and **SOF** over its serial versions.
- ▶ **p-apsp** and **SOF** improve as the number of threads increases.
- ▶ **p-apsp** achieved a **better speedup** with the increasing number of threads.

Table: Experiments with 300.000 ESTs with $\tau = 5$. The table shows the running time in seconds.

n. threads	apsp	p-apsp		SOF	
	time	time	speedup	time	speedup
1	397.17	463.33		82.80	
2		222.78	1.91	52.49	1.58
4		121.35	3.51	29.50	2.81
8		68.11	6.26	28.30	2.93
16		43.41	9.82	28.64	2.89
32		34.65	12.31	23.62	3.50

Experiments:

Peak Memory:

- ▶ The memory usage decreases as τ increases (number of overlaps).
- ▶ **SOF** uses **less memory** in all experiments.
- ▶ **p-apsp** memory usage is very similar to **apsp**, differing only by a constant factor.




Figure: Running time for varying values of τ (malloc_count library).

Outline

1. Introduction
2. Preliminaries
3. Related work
4. Parallel Algorithm
5. Experiments
6. Conclusion

Conclusion:

Future works:

- ▶ **Semi-external memory:** The GESA can be constructed in external memory and its blocks can be accessed as necessary, reducing the peak memory.
- ▶ **Different architecture model:** such as cloud distributed computing, possibly enabling the usage of hundreds of threads.

Thank you!



Gonnella, G. and Kurtz, S. (2012).

Readjoiner: a fast and memory efficient string graph-based sequence assembler.
BMC bioinformatics, 13(1):82.



Gusfield, D., M. Landau, G., and Schieber, B. (1992).

An efficient algorithm for the all pairs suffix-prefix problem.
Information Processing Letters, 41(4):181–185.



Ohlebusch, E. and Gog, S. (2010).

Efficient algorithms for the all-pairs suffix-prefix problem and the all-pairs substring-prefix problem.
Information Processing Letters, 110(3):123–128.



Rachid, M. H. and Malluhi, Q. (2015).

A practical and scalable tool to find overlaps between sequences.
BioMed Research International, 2015:1–12.



Simpson, J. T. and Durbin, R. (2010).

Efficient construction of an assembly string graph using the FM-index.
Bioinformatics, 26(12):i367–i373.



Tustumi, W. H., Gog, S., Telles, G. P., and Louza, F. A. (2016).

An improved algorithm for the all-pairs suffix-prefix problem.
Journal of Discrete Algorithms.